

# Table of Contents

1. Project Setup and Initialization
2. Data Loading and Initial Exploration
3. Target Label Analysis
4. Feature Selection by Missing Values
5. Feature Engineering and Preprocessing
  - 5.1. Separating Feature Types
  - 5.2. Creating Preprocessing Pipelines
6. Data splitting & Imputation
  - 6.1. Data Splitting and Imputation
  - 6.2.
7. Handling Class Imbalance with SMOTE
  - 7.1. Applying Synthetic Minority Over-sampling Technique (SMOTE)
  - 7.2. Training and Evaluation with SMOTE
8. Advanced Modeling and Hyperparameter Tuning (CatBoost)
  - 8.1. Hyperparameter Search with RandomizedSearchCV
  - 8.2. CatBoost, XgBoost, RF Model Evaluation
  - 8.3. Precision vs. Threshold Analysis
9. Results

---

## 1. Project Setup and Initialization

This initial step configures the computing environment and imports necessary libraries.

**Library Imports:** The necessary Python libraries are imported: `os` and `pandas` for data manipulation.

**Define File Paths:** Local paths are defined for the main training data (`train (6).csv`) and the churn label file (`train_churn_labels.csv`).

## 2. Data Loading and Initial Exploration

The primary feature dataset is loaded and summarized to understand its structure and statistical properties.

**Load Data:** The `train (6).csv` file is read into a pandas DataFrame named `data`.

**Descriptive Statistics:** The `.describe()` method is executed on the DataFrame, showing count, mean, standard deviation, and quartile information for the numeric columns. This summary is also saved to a file named `summary_features.csv`. The data contains 192 columns in the summary, with varying counts of non-missing values.

## 3. Target Label Analysis

The dataset containing the target variable is loaded to examine the class distribution.

**Load Labels:** The `train_churn_labels.csv` file is loaded into a DataFrame named `labels`.

**Label Values:** The unique values in the `Label` column are identified as -1 and 1.

**Class Imbalance Check:** The value counts for the target variable show a significant class imbalance: 46,328 instances with label -1 and only 3,672 instances with label 1.

## 4. Feature Selection by Missing Values

Columns are filtered based on the percentage of missing values to remove those with excessive sparsity.

**Missing Value Calculation:** The percentage of missing values for every column is calculated.

**Filtering Criterion:** Columns with less than 30% missing values are selected for the downstream analysis.

**Result:** This filtering step results in 67 columns being retained, creating a new DataFrame called `data_filtered_missing`.

## 5. Feature Engineering and Preprocessing

The selected features are prepared for model training by separating them by type and establishing a preprocessing pipeline.

### 5.1. Separating Feature Types

The 67 retained columns are separated into categorical (`object`, `category` dtypes) and numerical (`int64`, `float64` dtypes) feature lists.

- **Categorical features:** Variables like `Var192` through `Var228`.
- **Numerical features:** Variables like `Var6`, `Var7`, `Var13`, `Var21`, etc.

### 5.2. Creating Preprocessing Pipelines

- **Missing Indicator:** Flags columns with missing values.
- **Numerical Pipeline:**
  - Imputes missing values with the mean.
  - Applies `MinMaxScaler` for feature scaling.
- **Categorical Pipeline:**
  - Imputes missing values with the mean
  - Applies `labelEncoder` for categorical encoding.
- **Column Transformer:** Applies the numerical pipeline to numerical features and the categorical pipeline to categorical features.

## 6. Model Training and Evaluation

- Data split into training and testing sets (70/30 split) using `train_test_split`.
- Preliminary Pipeline combines preprocessing steps with a Logistic Regression model (`max_iter=5000`).
- Pipeline is fit on the training data (`X_train, y_train`).

## 6.2. Model Evaluation

- Generates probability predictions (`y_proba_lr`) on the test set.
- Classification threshold of 0.1 is applied to convert probabilities into binary predictions (`y_pred_lr`).
- Evaluated using a Classification Report and ROC-AUC score.

# 7. Handling Class Imbalance with SMOTE

To address the severe class imbalance identified in Step 3, the SMOTE technique is applied.

## 7.1. Applying Synthetic Minority Over-sampling Technique (SMOTE)

- SMOTE object is initialized and applied to the training data (`X_train, y_train`) to generate synthetic samples for the minority class (Label 1).
- Result is a new, balanced dataset (`X_resampled, y_resampled`).

## 7.2. Training and Evaluation with SMOTE

- Second Logistic Regression pipeline is trained using the SMOTE-resampled training data.
- Same threshold of 0.1 used for evaluation on the test set.

- Performance evaluated using Classification Report and ROC-AUC score, compared against the baseline model.

## 8. Advanced Modeling and Hyperparameter Tuning

An advanced gradient boosting model is introduced, and its hyperparameters are optimized.

### 8.1. Hyperparameter Search with RandomizedSearchCV

- Final pipeline constructed using the `ColumnTransformer` (preprocessing) and a CatBoost Classifier.
- Parameter grid (`param_grid_cat`) defined for tuning the CatBoost model.
- `RandomizedSearchCV` used to efficiently search the hyperparameter space (`n_iter=10, cv=3`).
- Search fitted on the original training data.

### 8.2. Model Evaluation

- Best estimator (`best_cat`) from randomized search is used to make probability predictions on the test set.
- Same threshold (0.1) applied to generate binary predictions (`y_pred_cat`).
- Best hyperparameters printed, performance presented via Classification Report and ROC-AUC score.

### 8.3. Precision vs. Threshold Analysis

- Precision-Recall Curve calculated using predicted probabilities and true test labels.
- Line plot visualizes the relationship between Precision and classification Threshold, allowing optimal threshold selection beyond the fixed 0.1 used in evaluations.
- 

## 9. Results

Fitting 3 folds for each of 15 candidates, totalling 45 fits

XGBoost Evaluation (Threshold=0.65):

Best XGB Params: {'subsample': 0.7, 'reg\_lambda': 5, 'reg\_alpha': 1, 'n\_estimators': 500, 'max\_depth': 4, 'learning\_rate': 0.01, 'colsample\_bytree': 0.8}

	precision	recall	f1-score	support
0	0.94	0.92	0.93	9266
1	0.20	0.24	0.22	734
accuracy			0.87	10000
macro avg	0.57	0.58	0.57	10000
weighted avg	0.88	0.87	0.88	10000

ROC-AUC: 0.7048997212862823

Fitting 3 folds for each of 10 candidates, totalling 30 fits

Random Forest Evaluation (Threshold=0.65):

Best RF Params: {'n\_estimators': 200, 'max\_features': 'log2', 'max\_depth': None, 'class\_weight': 'balanced'}

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

	0	0.93	1.00	0.96	9266
	1	0.50	0.00	0.01	734
accuracy				0.93	10000
macro avg		0.71	0.50	0.48	10000
weighted avg		0.90	0.93	0.89	10000

ROC-AUC: 0.666311633577622

Fitting 3 folds for each of 10 candidates, totalling 30 fits

`c:\users\rajud\anaconda3\lib\site-packages\catboost\core.py:1411:`

FutureWarning: iteritems is deprecated and will be removed in a future version.  
Use .items instead.

```
self._init_pool(data, label, cat_features, text_features, embedding_features,
embedding_features_data, pairs, weight,
```

CatBoost Evaluation (Threshold=0.65):

Best CatBoost Params: {'learning\_rate': 0.01, 'l2\_leaf\_reg': 7, 'iterations': 600, 'depth': 5}

		precision	recall	f1-score	support
	0	0.94	0.93	0.93	9266
	1	0.20	0.23	0.21	734
accuracy				0.88	10000
macro avg		0.57	0.58	0.57	10000
weighted avg		0.88	0.88	0.88	10000

ROC-AUC: 0.7024428178139175